

This is us: making CSAFE stronger each
week

CSAFE

2019-09-05

Contents

1	Prerequisites	5
2	Introduction	7
3	Project CC: Bullets and Cartridge Cases	9
3.1	Data Collection	10
3.2	Computational Tools	13
3.3	Similarity Scores	14
3.4	Analysis of Results	21
3.5	Communication of Results and Methods	21
3.6	People involved	21
4	Handwriting	23
5	Glass	25
6	Shoes	27
6.1	Longitudinal Shoe Study	27
6.2	Passive Shoe Recognition	29
6.3	Maximum Clique Matching	29
6.4	Cocoa Powder Citizen Science	29
7	Theoretical foundations	31
7.1	Nate's Updates 9/5/2019	31
8	Outreach activities	33

Chapter 1

Prerequisites

This is a *sample* book written in **Markdown**. You can use anything that Pandoc's Markdown supports, e.g., a math equation $a^2 + b^2 = c^2$.

The **bookdown** package can be installed from CRAN or Github:

```
install.packages("bookdown")  
# or the development version  
# devtools::install_github("rstudio/bookdown")
```

Remember each Rmd file contains one and only one chapter, and a chapter is defined by the first-level heading #.

To compile this example to PDF, you need XeLaTeX. You are recommended to install TinyTeX (which includes XeLaTeX): <https://yihui.name/tinytex/>.

Chapter 2

Introduction

This section will become the section for the administrative updates/organization once we have figured out how to use all of the bookdown features for our purposes.

You can label chapter and section titles using `{#label}` after them, e.g., we can reference Chapter 2. If you do not manually label them, there will be automatic labels anyway, e.g., Chapter 5.

Figures and tables with captions will be placed in `figure` and `table` environments, respectively.

```
par(mar = c(4, 4, .1, .1))  
plot(pressure, type = 'b', pch = 19)
```

Reference a figure by its code chunk label with the `fig:` prefix, e.g., see Figure 2.1. Similarly, you can reference tables generated from `knitr::kable()`, e.g., see Table 2.1.

```
knitr::kable(  
  head(iris, 20), caption = 'Here is a nice table!',  
  booktabs = TRUE  
)
```

You can write citations, too. For example, we are using the **bookdown** package (Xie, 2019) in this sample book, which was built on top of R Markdown and **knitr** (Xie, 2015).




Figure 2.1: Here is a nice figure!

Table 2.1: Here is a nice table!

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.0	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.0	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa
4.6	3.4	1.4	0.3	setosa
5.0	3.4	1.5	0.2	setosa
4.4	2.9	1.4	0.2	setosa
4.9	3.1	1.5	0.1	setosa
5.4	3.7	1.5	0.2	setosa
4.8	3.4	1.6	0.2	setosa
4.8	3.0	1.4	0.1	setosa
4.3	3.0	1.1	0.1	setosa
5.8	4.0	1.2	0.2	setosa
5.7	4.4	1.5	0.4	setosa
5.4	3.9	1.3	0.4	setosa
5.1	3.5	1.4	0.3	setosa
5.7	3.8	1.7	0.3	setosa
5.1	3.8	1.5	0.3	setosa

Chapter 3

Project CC: Bullets and Cartridge Cases

For both bullets and cartridge cases we are dealing with several inter-related aspects, that we want to address independently.

Those are:

1. data collection
2. computational tools
3. similarity scores
 1. for bullet lands:
 - a. crosscut identification
 - b. groove location
 - c. curvature removal
 - d. alignment of signatures
 - e. feature extraction
 - f. matching with trained Random Forest
 2. for breech faces
4. analysis of results
5. communication of results and methods

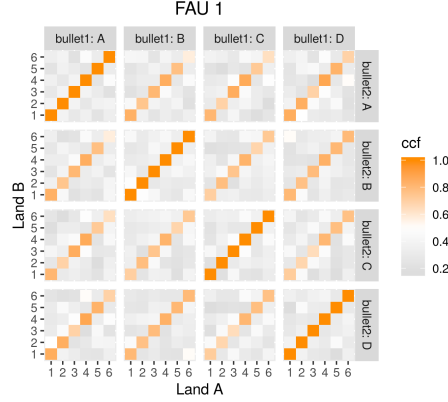


Figure 3.1: Results from assessing scans of barrel FAU 1 similarity.

3.1 Data Collection

3.1.1 LAPD

All bullets are collected by Srinivasan Rathinam, LAPD.

3.1.1.1 Main study

4 bullets per barrel for 626 Beretta 92 F/FS firearms , ammunition used are 9 mm Luger Winchester 115 grain with a Copper surface.

scans are on Raven.

evaluation: Yawei is going to work through all 626 barrels of knowns to assess similarity scores

3.1.1.2 follow-up study

4 bullets per barrel for 96 of the original 626 Beretta firearms using different ammunition

bullets are being scanned




Figure 3.2: Results from assessing scans of barrel FAU 2 similarity.

3.1.2 Hamby Sets

Scans for Hamby Sets 10, 36, 44, and 224

Scans for 3 replicates of clones for Hamby 224

3.1.3 Houston Tests

contact: Melissa Nally, Houston FSI

3.1.3.1 Pre-study

3 kits with 23 bullets each

evaluation included in submission to JFI

3.1.3.2 Study

4 kits with 20 bullets each

scans done, evaluation finished, some scans of doubtful quality




Figure 3.3: Bullet-to-bullet similarity scores for questioned bullets (y-axis) compared to all other bullets of the test set (x-axis).

3.1.4 Houston Persistence

contact: Melissa Nally, Houston FSI

8 barrels with 40 fired bullets each

3.1.5 St Louis persistence

contact: Steve Kramer, St Louis PD

2 barrels with 192 fired bullets each (2 bullets collected every 25 shots)

3.1.6 DFSC Cartridge cases

Breech face data for knowns are scanned and available on a private github repository

evaluation

3.2 Computational Tools

3.2.1 x3ptools

`x3ptools` is an R package for working with files in x3p format. x3p is an ISO standard for describing 3d topographic surface measurements. `x3ptools` is available on CRAN, i.e. can be installed with the command `install.packages("x3ptools")`. The development version is available from github. Installation instructions and basic usage can be found at <https://heike.github.io/x3ptools/>

3.2.2 bulletxtrctr

`bulletxtrctr` is a developmental R package available from github (see <https://heike.github.io/bulletxtrctr/>) that allows an assessment of similarity scores using the data extraction pipeline described in Hare et al. (2016).

3.2.3 grooveFinder

`grooveFinder` is a developmental R package providing different methods for identifying the location of grooves in scans of bullets. Installation instructions and some basic usage can be found at <https://heike.github.io/grooveFinder/>




Figure 3.4: 2-dimensional visualization of example bullet br411 with .999 strength threshold




Figure 3.5: 3-dimensional visualization of example bullet br411 with .999 strength threshold

3.3 Similarity Scores

3.3.1 Bullet Lands

3.3.1.1 Approaches to identify groove locations

3.3.1.1.1 Hough Transform Method for Identifying Grooves

Charlotte 9/5/19 Update:

Will fill in with more detail later

Current Goals: - Iron-out issues with consistency of units with `get_hough_grooves`.

I believe there are some issues translating from the 2-d visualization to the 3-d visualization that might have to do with inconsistent unit inputs? For Example

So either somethin is wrong with `get_mask_hough` or something is funky with the units.

- Also need to think of including a sort of rounding component where lines with slopes that are practically infinite can be viewed as a vertical line
- Compare Hough results with manual identification using score calculations from Kiegan.
- Write up results in Hough Groove Paper (It's coming I promise)




Figure 3.6: Comparing a cell in image 1 to a larger region in image 2. We wish to find the translations of the image 1 cell that yield the highest correlation within the image 2 region.

- Create graphical images to explain line selection method
- Include 2-d and 3-d visualizations of Hough groove area identifications
- Include crosscut visualization and comparison in results

3.3.2 Cartridge Cases

3.3.2.1 Congruent Matching Cells (CMC) algorithm for comparing cartridge case breech faces

Joe 9/5/19 Update: Dealing with missing values in the x3p scans continues to be an issue. The Fast Fourier Transform method for calculating cross-correlation can't handle missing data in an image, so we've attempted a few "fixes" that haven't necessarily turned out as well as expected. One idea we had was to replace the NA values in a cell with the average pixel value. However, this is artificially introducing a signal where before there was none. This can (and demonstrably has) led to inflated/incorrect correlations between cells that shouldn't have much at all in common. Unfortunately, this may be the only solution if we still wish to adhere to the CMC algorithm as described in Song et al. (2015). One improvement that I've implemented is to "crop out" the rows and columns of an image that only contain NAs. This at least means that we've weakened the strength of the artificial signal relative to the breechface's signal.

Below is a series of images that illustrate how we might compare a cell in one image to a region of another image.

For the sake of an example, let's focus on the blue outlined cell in image 1. Our goal is to use the image 1 cell to "search" a corresponding larger region in image 2 for the horizontal/vertical translations needed to produce the highest correlation. Below is a zoomed-in version of the blue outlined image 1 cell on the left and the larger image 2 region (approximately: I made the gridded image above by-hand outside of R while the images below are from R). The image 1 cell may look larger than the image 2 region, but we can see from the axes that

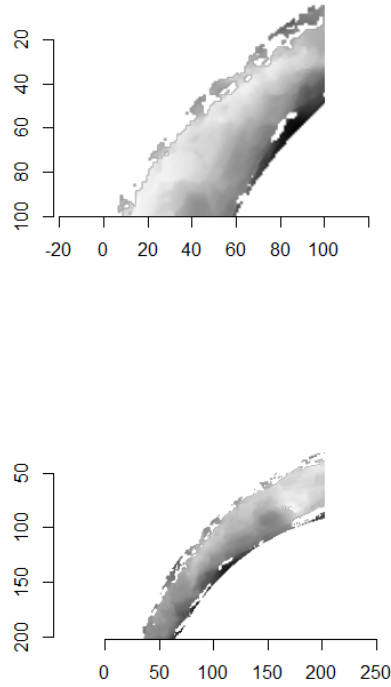


Figure 3.7: (Left) A cell from image 1. (Right) A region from image 2 centered in the same location as the image 1 cell, yet quadruple the area.

the image 2 region is indeed larger. Any white pixels in the two images are NA values that need to be dealt with in some way before we can use FFTs to calculate the cross-correlation.

As already discussed above, one “solution” is to replace the NA values with the average pixel value of each image. However, to avoid creating a stronger artificial signal than necessary, we can crop-out the NA rows and columns from the two images above. Below is the cropped version of the two images. The cropping doesn’t produce significantly different images in this case, but you could imagine other examples in which a cell has captured only small amount of breechface in the corner. Such examples are fairly common and cropping significantly changes the resulting correlation values.

The last step before calculating correlation for these cells is to replace the remaining NAs with the average pixel value. This is shown below.




Figure 3.8: The same images as above after cropping NA rows/columns.

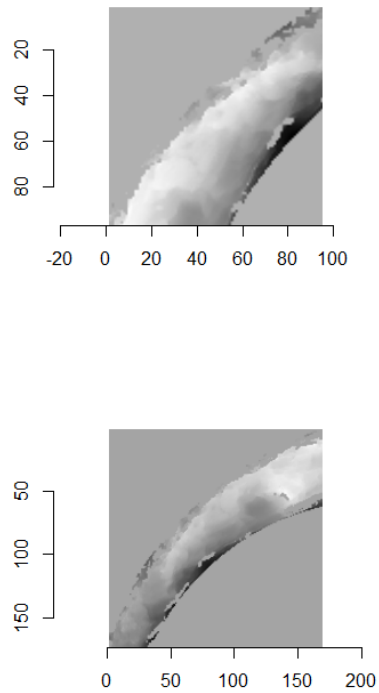


Figure 3.9: The NA-cropped images with remaining NAs replaced with the image's average pixel values.

The cross-correlation is then calculated between these two images via a standard fast fourier transform process (see Cross-Correlation Theorem). The benefit of using such a process is that (as the name suggests) it's faster than calculating the raw correlation between the two images. Also, the translations that produce the highest correlation between the image 1 cell and the image 2 region fall out of the calculation for free.

This pre-processing/cross-correlation calculation procedure is repeated for every cell in image 1 that contains breech face. Because it is not valid to assume that the two images are rotationally aligned by default, we perform the same procedure repeatedly while rotating image 2. Currently, we perform a “rough” grid search of $\theta \in [-177.5, 180]$ by increments of 2.5° . Theoretically, the final results tell us how we need to horizontally/vertically translate and rotate the two images to be correctly aligned.

3.3.2.2 Congruent Matching Tori: a promising solution to the missing value problem

As discussed above, dealing with missing values is proven to be a pain. The good news is that the currently-implemented CMC as described above yields results very similar to those published in Song et al. (2015) that originally describes that CMC algorithm. While our results seem to agree with currently published results, it would be nice if we could avoid needing to artificially replace missing values. We can do so if, rather than breaking up the circular breech face scans into disjoint squares, we break up the breech face into donut-shaped regions containing only breech face. Below is an example of such a toroidal region.

By comparing such regions instead of the square cells, we would presumably only need to fill in a few missing value “holes” in the breech face scan rather than completely replacing a non-existent signal with an artificial one. In the near-future, I hope to finish up the pre-processing needed for this Congruent Matching Tori method by performing a polar transformation on these images to make them into strips that can easily be compared via an FFT.




Figure 3.10: (Left) The original breech face scan image. (Right) A donut-shaped region cut out of the original image.

3.4 Analysis of Results

3.5 Communication of Results and Methods

3.6 People involved

3.6.1 Faculty

- Heike Hofmann
- Susan VanderPlas

3.6.2 Graduate Students

- Ganesh Krishnan
- Kiegan Rice
- Nate Garton
- Charlotte Roigers
- Joe Zemmels
- Yawei Ge

3.6.3 Undergraduates

- Talen Fisher (fix3p)
- Andrew Maloney
- Mya Fisher, Allison Mark, Connor Hergenreter, Carley McConnell, Anyesha Ray (scanner)

Chapter 4

Handwriting

We describe our methods for going about the handwriting project here.

Chapter 5

Glass

Chapter 6

Shoes

6.1 Longitudinal Shoe Study

Github repository

6.1.1 Paper describing the database

Paper subdirectory of Github repository

Goal:

- Describe experiment
- Describe database function
- Publicize data for analysis by others in the community

Methods and Data Description

Methods and data description handed off to Alicia for editing

Data Analysis Tools

- Working with the `EBImage` package - very fast processing of images

6.1.1.0.1 Film and Powder Images

Analysis Steps:

1. Create threshold mask

- a. Blur image (circular/gaussian blur, diameter 15)
 - b. Invert the image
 - c. Threshold image (adaptive threshold, 10 x 10 region, keep anything with an average higher than 0.025 from the mean)
 - d. Create mask
Default parameters selected by visually screening several shoes:

(default parameters $\text{rad1} = 5$, $\text{rad2} = 151$, $\text{proportion} = 1.5 \times \text{area of rad2 in px} / \text{area of image in px}$)
 1. erode mask image (circle, diameter rad1)
 2. dilate mask image (circle, diameter rad2)
 3. label disjoint regions of the image
 4. prune small image regions ($\text{area} < \text{proportion parameter}$)
 - e. Fill in mask holes
 - f. Expand mask to capture entire shoe region
 1. set background color
 2. create dataframe of useful (non-background) pixels
 3. fill in holes and concave regions in mask, then expand by `expand_rad` vertically and horizontally (similar to “convex hull”, but with additional expansion radius)
2. Mask image to remove extra variability unrelated to the shoe
 3. Threshold masked image?
Con: Lose grey information; Pro: fully remove background
 4. Compromise: Keep grey pixels from thresholded, masked image (e.g. use 3. as a mask), then renormalize

6.1.1.0.2 Wear Characterization

Ideas:

- average intensity of cleaned image
- length of border/edges detected

6.2 Passive Shoe Recognition

CoNNOR: Convolutional Neural Network for Outsole Recognition

- Label images of shoes according to geometric classification scheme
- Use convolutional base of pretrained CNN VGG16 and train a new classifier on labeled features

Github repository for CC on CoNNOR

Github repository for paper submitted to Forensic Science International

Exploring new directions:

- Truncate convolutional base and train random forest on features
 - Could replace fully connected layers of neural net as classifier
 - Importance score can filter/reduce the number of features
 - *Currently training on HPC, 3 days in*
- Fully convolutional networks (FCNs)
 - Unsupervised segmentation to assess current classification scheme
 - Handle whole shoe image of any size (instead of only 256x256 pixel images)

References for CNNs and FCNs

Stack Exchange post explaining patchwise training

“Learning Hierarchical Features for Scene Labeling”: describes an application of multi-scale CNNs and image pyramids

“Pyramid methods in image processing”: classic paper from 1984 explaining pyramid methods

“Fully Convolutional Networks for Semantic Segmentation”

“W-Net: A Deep Model for Fully Unsupervised Image Segmentation”

6.3 Maximum Clique Matching

6.4 Cocoa Powder Citizen Science

Chapter 7

Theoretical foundations

7.1 Nate's Updates 9/5/2019

- Currently in Virginia
- RA for this semester (year?) is under Danica
- **Central Goals:**
 - continue work started by Danica and Peter Vergeer on the analysis of likelihood ratios
 - study the differences between specific source (SS) and common source (CS) likelihood ratios (LRs) in an information theoretic way
 - does the CS or SS LR have more “information”?
 - can be the CS or SS hypotheses (prosecution or defense) be formally compared in terms of being easier to “prove” or “disprove”?
- **Basic Setup**
 - H_p , H_d are CS prosecution and defense hypotheses upon which we will place priors
 - A and B are discrete r.v.'s representing two “sources” of evidence
 - * distributions for A and B defined conditionally based on the hypothesis
 - * SS hypothesis is represented by the conditional random variable $H_p|A$
 - X is data coming from A , Y is data coming from B
 - compare information contained in (X, Y) about H_p and $H_p|A$
 - * this is what I'm working on now
 - * important quantities:
 - Kullback-Leibler divergences, entropy, mutual information

Chapter 8

Outreach activities

Bibliography

- Hare, E., Hofmann, H., and Carriquiry, A. (2016). Automatic Matching of Bullet Lands. *ArXiv e-prints*.
- Xie, Y. (2015). *Dynamic Documents with R and knitr*. Chapman and Hall/CRC, Boca Raton, Florida, 2nd edition. ISBN 978-1498716963.
- Xie, Y. (2019). *bookdown: Authoring Books and Technical Documents with R Markdown*. R package version 0.13.